

Uninformed Search

CSE 4617: Artificial Intelligence



Isham Tashdeed
Lecturer, CSE



Machines that are **Intelligent**

- Some have defined intelligence through the lens of **rationality** while others see it compared with **human** performance
- Another way to categorizing intelligence would be to consider intelligence being a property of **internal thought process** while others classify intelligence on **behavior**

There are *four* combinations of these schools of thought.


Think Rationally

- Rationally → Doing the *right* thing
- What are the rules that govern correct thought?
- Let us infuse those rules into machines




Think like people

- Reverse engineering the brain
- What are the processes going on inside our head
- Cognitive science




Act like people

- The Turing Test approach
- Almost all of modern Deep Learning advancements
- Fooling people by acting like them
- Tells us more about people than intelligence



Act rationally

- Focusing on the decisions made by the system
- Rational \rightarrow Optimal
- System *acts* as to achieve the best possible outcome



“Big Data”, “AI powered”,
“Prompting”, “Agentic
Approach”



Just behaves
rationally



Rational Decisions

- Maximally achieve a pre-defined goal
- We only care about the actions → Not the thoughts / rationale behind those actions

Maximize your expected utility

Rational Agents

- Agent → An entity that *perceives* and *acts*
- Rational Agent → Selects the actions that would maximize the utility



Reflex Agents

- Choses actions based on only current perception / memory
- Does not consider the future consequences of its action
- May have a memory / model of the world
- Only considers how the world is currently, rather than what it could be after the action
- Pros → Very quick

But is it **rational**?


Reflex Agents

- Choses actions based on only current perception / memory
- Does not consider the future consequences of its action
- May have a memory / model of the world
- Only considers how the world is currently, rather than what it could be after the action
- Pros → Very quick


But is it **rational**?

- We only consider the outcome of the action, so it can be rational

Reflex Agents



Reflex Agents




Planning Agents

- Takes action based on consequences (hypothesized) of the action
- **Must** have a model of the world. The model has to be updated as a response to the action taken
- Must formulate a goal
- A planning agent must be:
 - Optimal
 - Complete




Planning Agents



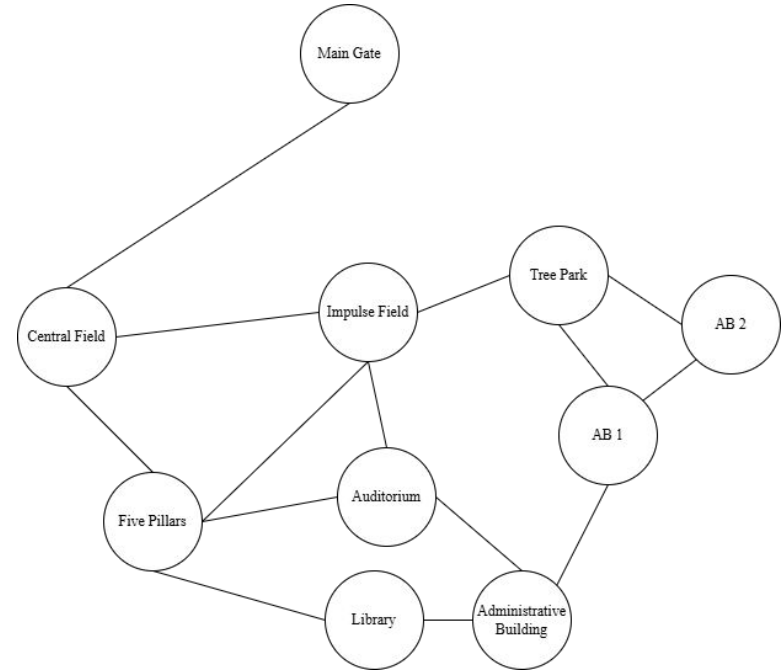
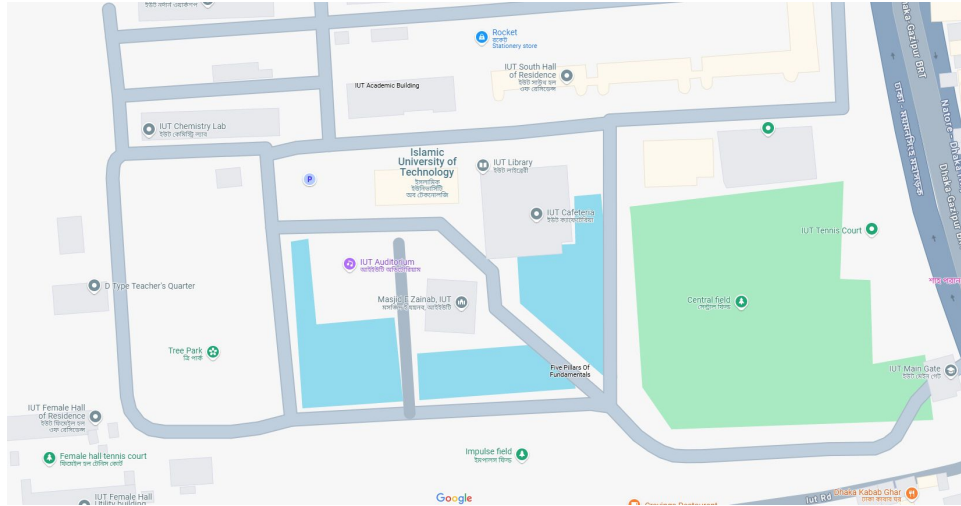
Search Problems

Defining a Search Problem

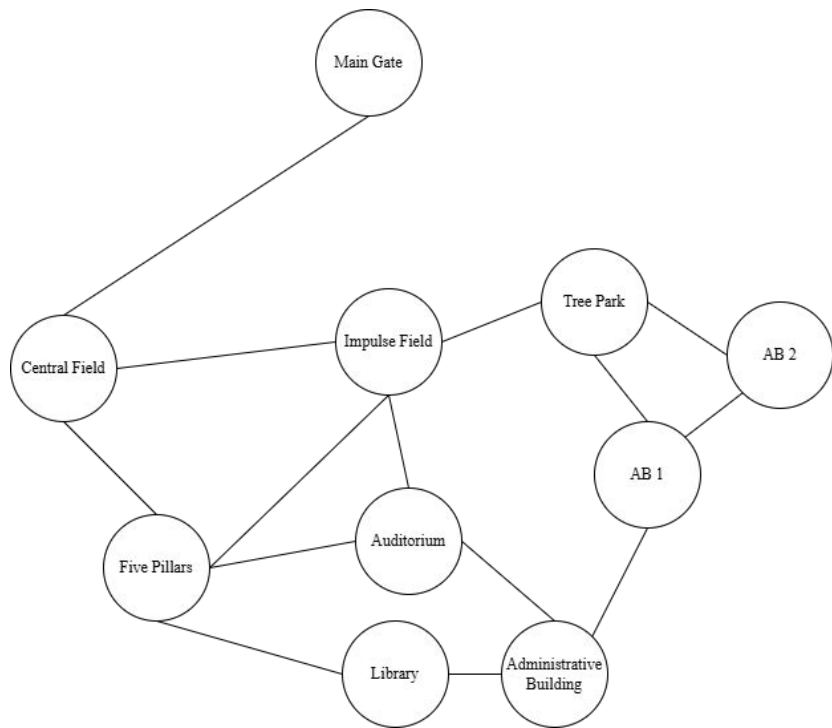
- **State space** → Set of all possible scenarios
 - A singular state defines the agent's condition at the current time
 - Actions can take an agent from one state to another
- **Successor function** → Provides consequent states
 - Takes an action and provides possible states from current state
 - May take *cost* into consideration (if applicable)
- **Start state**
- **Goal test**
- **Solution** → Sequence of actions, starting from start state to goal state



Example: Getting around IUT




Example: Getting around IUT




- State space?
 - Locations
- Successor Function?
 - Connections / Roads
- Start state?
 - Main gate
- Goal test?
 - Current state == AB2 ?
- Solution?
 - Path from start to goal state

Example: Maze Traversal



- State space?
- Successor Function?
- Start state?
- Goal test?
- Solution?


Example: Maze Traversal



- State space?
 - (x, y) location
- Actions: NSEW
- Successor Function?
 - Update the location of the agent
- Start state?
- Goal test?
 - $\text{Current } (x, y) == \text{END?}$

Only have to think about solving the maze


Example: Eat All Food



- State space?
 - (x, y) location, Food locations
- Actions: NSEW
- Successor Function?
 - Update the location and food
- Start state?
- Goal test?
 - All food eaten?

Have to think about solving the maze and eating all the food pellets


Size of the State Space



- Agent positions: 120
- # of food: 30
- # of ghosts: 2
- Ghost positions: 12
- Agent can face: NSEW

Total World states?


Size of the State Space



- Agent positions: 120
- # of food: 30
- # of ghosts: 2
- Ghost positions: 12
- Agent can face: NSEW

Total World states? $\rightarrow 120 * 2^{30} * 12^2 * 4$

Size of the State Space



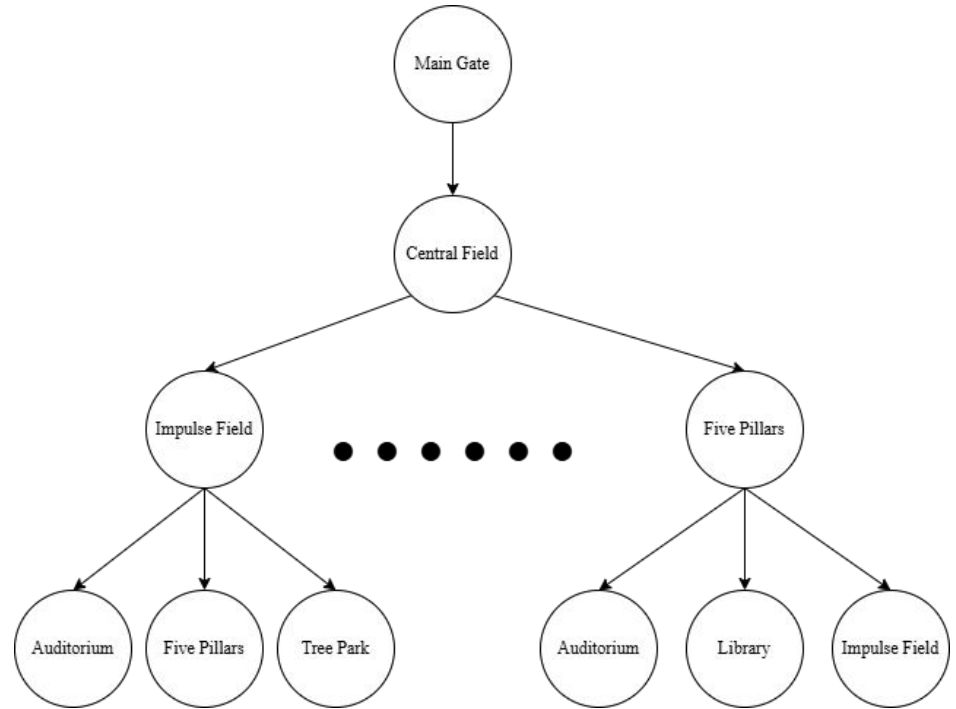
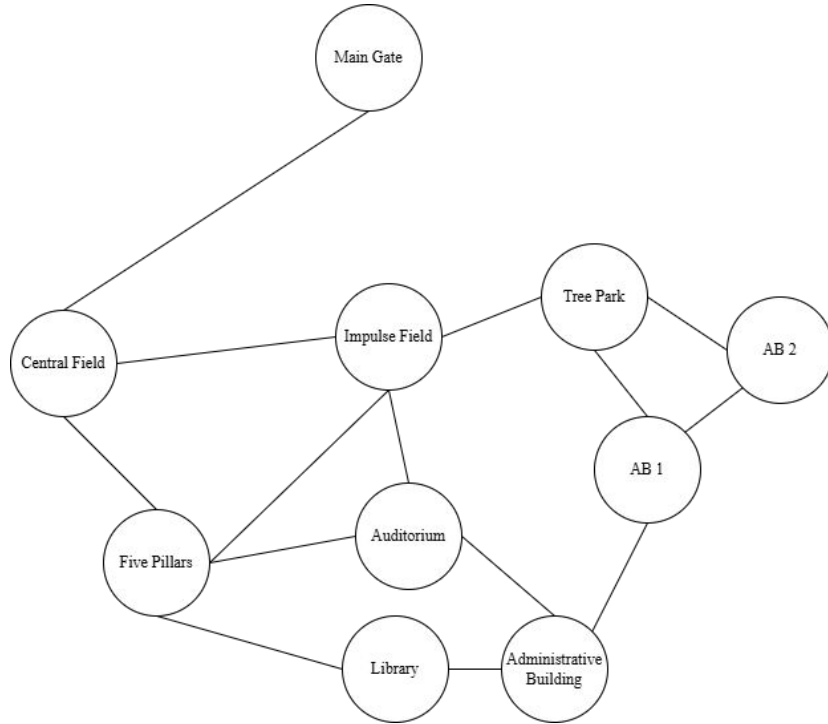
- Agent positions: 120
- # of food: 30
- # of ghosts: 2
- Ghost positions: 12
- Agent can face: NSEW

Total World states? $\rightarrow 120 * 2^{30} * 12^2 * 4$

Maze traversal $\rightarrow 120$


Eat all food $\rightarrow 120 * 2^{30}$

State Space Graphs and Search Trees



How big is the Search Tree?


Consider this 4-state graph:



Let, s be the root.

How big is the Search Tree?

Consider this 4-state graph:

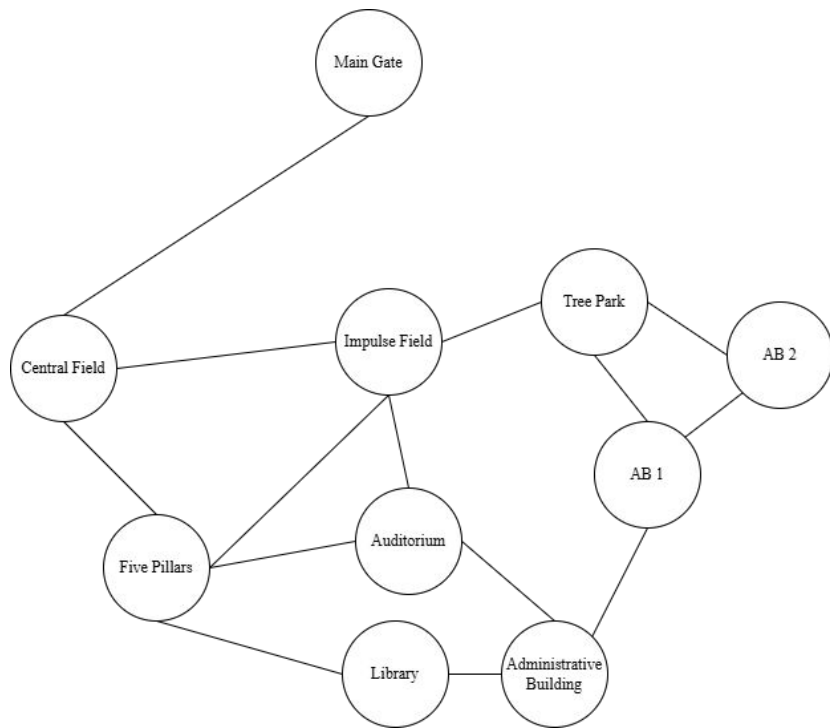


Let, s be the root.

- The search tree is infinite!
- Usually search trees have lots of repeated structures
- We can not explore the whole search tree at once

Tree Search

General Tree Search



- Expand child nodes
- Maintain a list/backlog of partial routes
- Expand as few nodes as possible


General Tree Search

```
TREE-SEARCH (problem, strategy) → returns a solution or failure
  initialize tree search using the initial state of the problem
  loop do:
    if there are no candidates for expansion:
      then return failure

    choose a leaf node according to strategy
    if chosen node contains a goal state:
      then return the solution
    else:
      Expand the node & add the resulting nodes to the search tree
  end
```

Depth-First Search

Depth-First Search



- Expand the deepest node first
- Fringe is a LIFO stack


Depth-First Search

```
DFS (problem) → returns a solution or failure
  initialize tree search using the initial state of the problem
  loop do:
    if there are no candidates for expansion:
      then return failure

    choose the most recently added leaf node
    if chosen node contains a goal state:
      then return the solution
    else:
      Expand the node & add the resulting nodes to the search tree
  end
```

Search Algorithm Properties

- Complete → Guaranteed to find a solution if it exists
- Optimal → Guaranteed to find the least cost path
- Time complexity → How many nodes does it expand?
- Space Complexity → How much space does the fringe take




Let's consider a simplified search tree

Search Algorithm Properties

- b is the branching factor
- m is the maximum possible depth
- The red nodes are the solutions at various depths


How many nodes in the whole tree?

$$1 + b + b^2 + b^3 + \dots + b^m = O(b^m)$$



DFS Properties


- Is it complete?
 - Only if m is finite
- Is it optimal?
 - No, only finds the left-most solution
- Time complexity?
 - At worst case, needs to process the whole tree
 - $O(b^m)$
- Space complexity?
 - Only has siblings on the path to root
 - $O(bm)$



Breadth-First Search

Breadth-First Search

Breadth
First
Search



- Expand the shallowest node first
- Fringe is a FIFO queue


Breadth-First Search

```
BFS (problem) → returns a solution or failure
  initialize tree search using the initial state of the problem
  loop do:
    if there are no candidates for expansion:
      then return failure

    choose the least recently added leaf node
    if chosen node contains a goal state:
      then return the solution
    else:
      Expand the node & add the resulting nodes to the search tree
  end
```

BFS Properties

- Is it complete?
 - s is finite if a solution exists
- Is it optimal?
 - Only if all the costs are equal
- Time complexity?
 - Process all nodes above the shallowest solution at depth s
 - $O(b^s)$
- Space complexity?
 - At worst case, needs to store all the nodes at depth s
 - $O(b^s)$



Which one is better?


Iterative Deepening

DFS (Space advantage) + BFS (Time advantage)

- Run DFS with depth limited to 1 layer...
- Run DFS with depth limited to 2 layers...
- Run DFS with depth limited to 3 layers...
- Until shallowest solution is found!

Is it redundant? → YES 😞

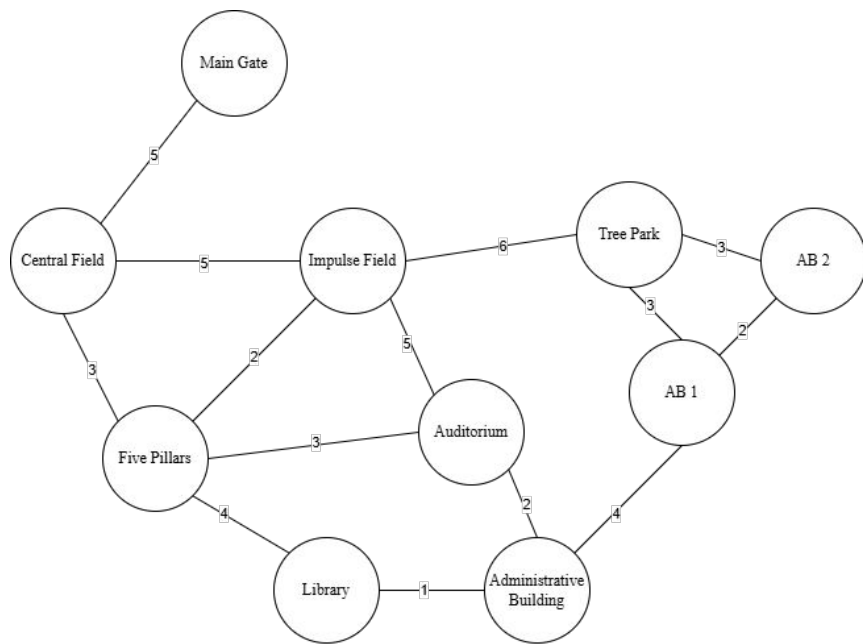
- Usually the number of nodes in the lower layers always dominates all other layers



MFW I explore the
deeper layers

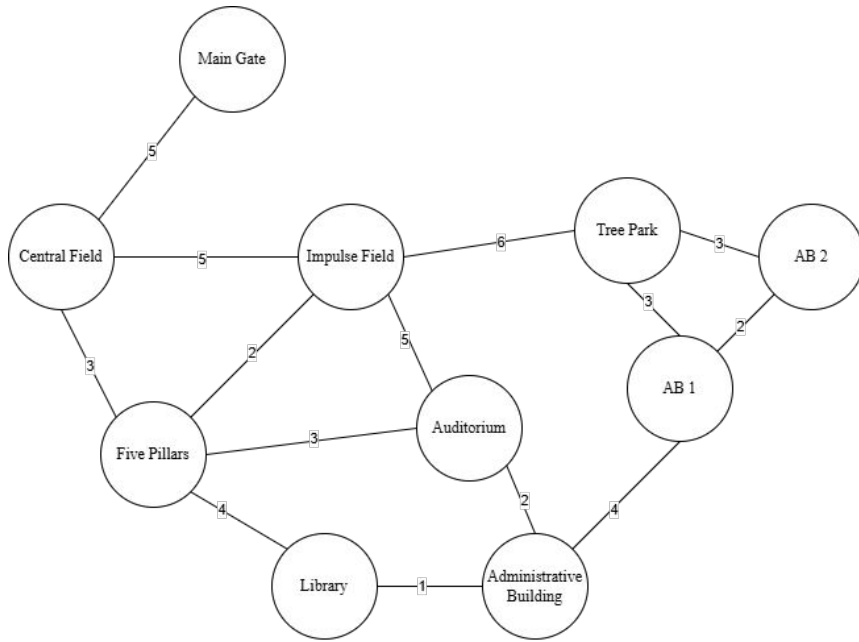
Cost Sensitive Search

Cost Sensitive Search



- What if there were some cost added to each path?
- BFS is not optimal as it won't find the least cost path, only the shallowest solution

Uniform Cost Search




- Expand the cheapest node first
- Fringe is a priority queue
 - Priority \rightarrow Cumulative Cost

Uniform Cost Search

```
UCS (problem) → returns a solution or failure
  initialize tree search using the initial state of the problem
  loop do:
    if there are no candidates for expansion:
      then return failure


    choose the leaf node with lowest cumulative cost
    if chosen node contains a goal state:
      then return the solution
    else:
      Expand the node &
      add the resulting nodes and their cumulative costs
      to the search tree
  end
```

Uniform Cost Search



UCS Properties


- Is it complete?
 - Only if best solution has a finite cost & the minimum cost is positive
- Is it optimal?
 - Yes. But why?
- Time complexity?
 - Process all nodes with cost less than the cheapest solution
 - If optimal solution costs C^* and each edge costs at least ϵ
 - Effective depth $\rightarrow O(b^{C^*/\epsilon})$
- Space complexity?
 - $O(b^{C^*/\epsilon})$



UCS Issues


- UCS explores in every “direction” as long as the costs are lower
- No information about where the goal state is located

How could we infuse the idea of where the goal is located?



Uninformed Search Algorithms

- DFS, BFS and UCS are all the same except for the fringe strategies
- Can implement any of these algorithms just by changing the priority function in a priority queue



Additional Resources

- [Maze Solving- Computerphile](#)
- [Theseus and the Minotaur | Exploring State Space](#)
- [Monte-Carlo Tree Search](#)

Thank you