

Reinforcement Learning II

CSE 4617: Artificial Intelligence




Isham Tashdeed
Lecturer, CSE



Formalizing RL Problems

- Still assume an MDP
 - A set of states $s \in S$
 - A set of actions per state A
 - A probability model $T(s, a, s')$
 - A reward function $R(s, a, s')$
- Still trying to find the most optimal policy π^*
- We **do not know** $T(s, a, s')$ and $R(s, a, s')$
- Must actually try out actions and states to learn



MDP vs RL

Known MDP: Offline Solution

Goal

Compute V^* , Q^* , π^*

Evaluate a fixed policy π

Technique

Value / policy iteration

Policy evaluation

Unknown MDP: Model-Based

Goal

Compute V^* , Q^* , π^*

Evaluate a fixed policy π

Technique

VI/PI on approx. MDP

PE on approx. MDP

Unknown MDP: Model-Free

Goal

Compute V^* , Q^* , π^*

Evaluate a fixed policy π

Technique

Q-learning

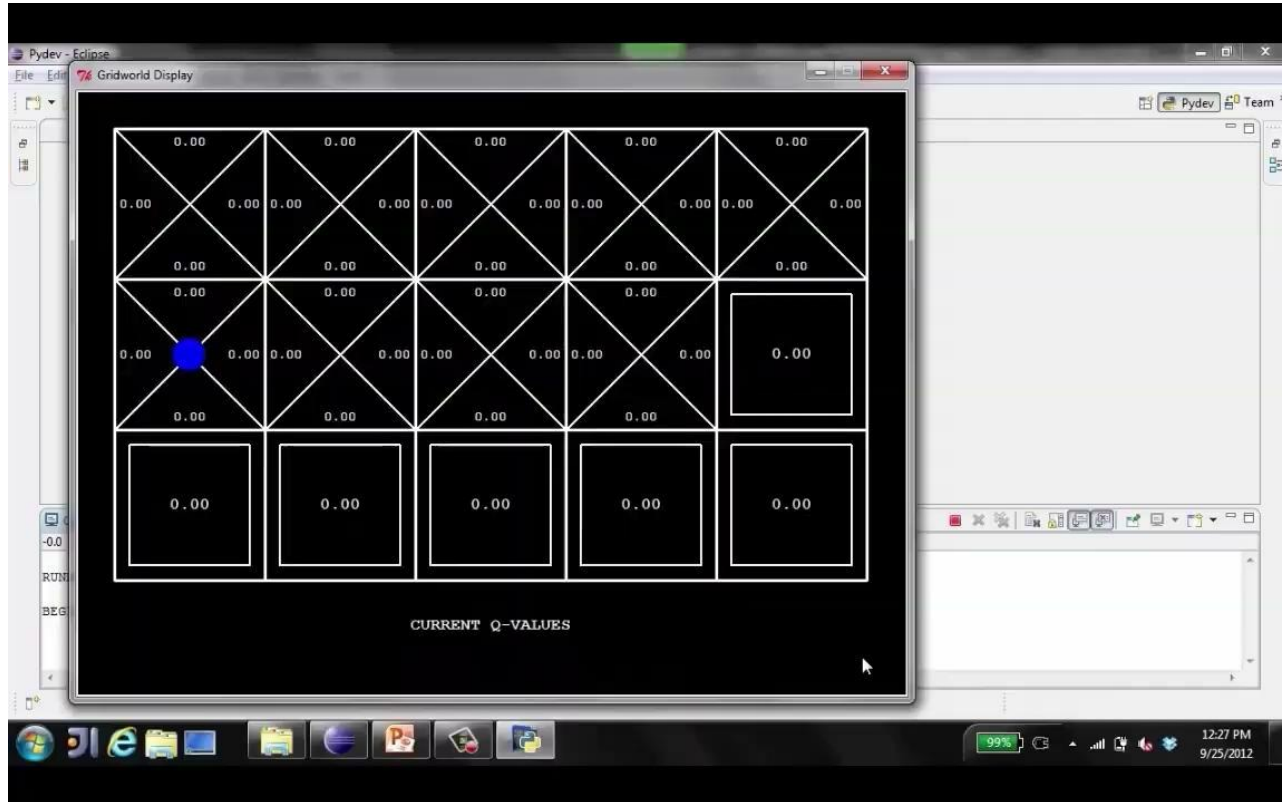
Value Learning

Q-Learning

Learn $Q(s, a)$ as you go

- Receive a sample (s, a, s', r)
- Consider the old estimate $Q(s, a)$
- New sample estimate $\rightarrow sample = R(s, a, s') + \gamma \max_{a'} Q(s', a')$
- Incorporate the new estimate into the running average
 - $Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha (sample)$

Q-Learning



Q-Learning Properties

- Q-learning converges to optimal policy \rightarrow even if you're acting suboptimally!
- This is called **off-policy learning**

Cons:

- You have to explore enough
- You have to eventually make the learning rate (α) small enough
- But not decrease the learning rate (α) too quickly

How to Explore?

Schemes for forcing exploration

- Random action (ϵ -greedy):
 - Every time step, flip a coin
 - With a very small probability (ϵ), stop following the established policy and act randomly
 - With a large probability ($1-\epsilon$), keep following the established policy
 - Setting ϵ to a fixed value ensures sufficient exploration but the agent will not get a chance to use what it has learned
 - Gradually decrease ϵ

Exploration Functions

Explore areas whose **badness/goodness** is not (yet) established, eventually stop exploring

- Takes a value estimate u and a visit count n , and returns an optimistic utility: $f(u, n) = u + k/(n + 1)$
- Modified Q-update:

$$Q(s, a) \leftarrow_{\alpha} R(s, a, s') + \gamma \max_{a'} f(Q(s', a'), N(s', a'))$$



Regret

Regret



- Even if you learn the optimal policy, you still make mistakes along the way!
- Regret is the difference between the total expected rewards and optimal expected rewards
- Gives a notion of how quickly an agent learns
- Minimizing regret means you will optimally learn how to be optimal!
- Random exploration and exploration functions both end up optimal, but random exploration has higher regret

Approximate Q-Learning

Generalizing Across States



- Simple Q-Learning keeps a table of all q-values → Or in a dictionary!
- In realistic situations, we cannot possibly learn about every single state! → Too many states!
 - Too many states to visit them all in training
 - Too many states to hold the q-tables in memory
- A better idea is to generalize to unseen states
 - Learn about a small number of states
 - Generalize the experience to novel yet similar states
 - This idea can be borrowed from machine learning

Generalizing Across States


Let's say we discover through experience that this state is bad:



In naïve q-learning, we know nothing about this state:



Or even this one!



Feature-based Representations



- Describe a state using a vector of features
- Features are functions from states to real numbers (often $[0, 1]$) that capture important properties of the state
- They can be hand crafted \rightarrow Similar to evaluation functions of a state
 - Distance to closest ghost
 - Distance to closest food
 - Number of ghosts
 - $1 / (\text{distance to food})^2$
 - Is Pacman in a tunnel? (0/1) \rightarrow True/False
- Or these vectors can be learned \rightarrow Deep Learning

Linear Value Functions

Using a feature representation, we can write a q-function (or value function) for any state using a few weights:

$$V(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

- Q-learning will learn the weights $(w_1, w_2 \dots w_n)$ to approximate the values of the states
- If the number of features are not adequate, your agent will not be able to tell states apart, even though they are very different from each other

Approximate Q-Learning

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

Q-Learning with linear Q-functions:

$$\text{transition} = (s, a, r, s')$$

$$\text{difference} = \left[r + \gamma \max_{a'} Q(s', a') \right] - Q(s, a)$$

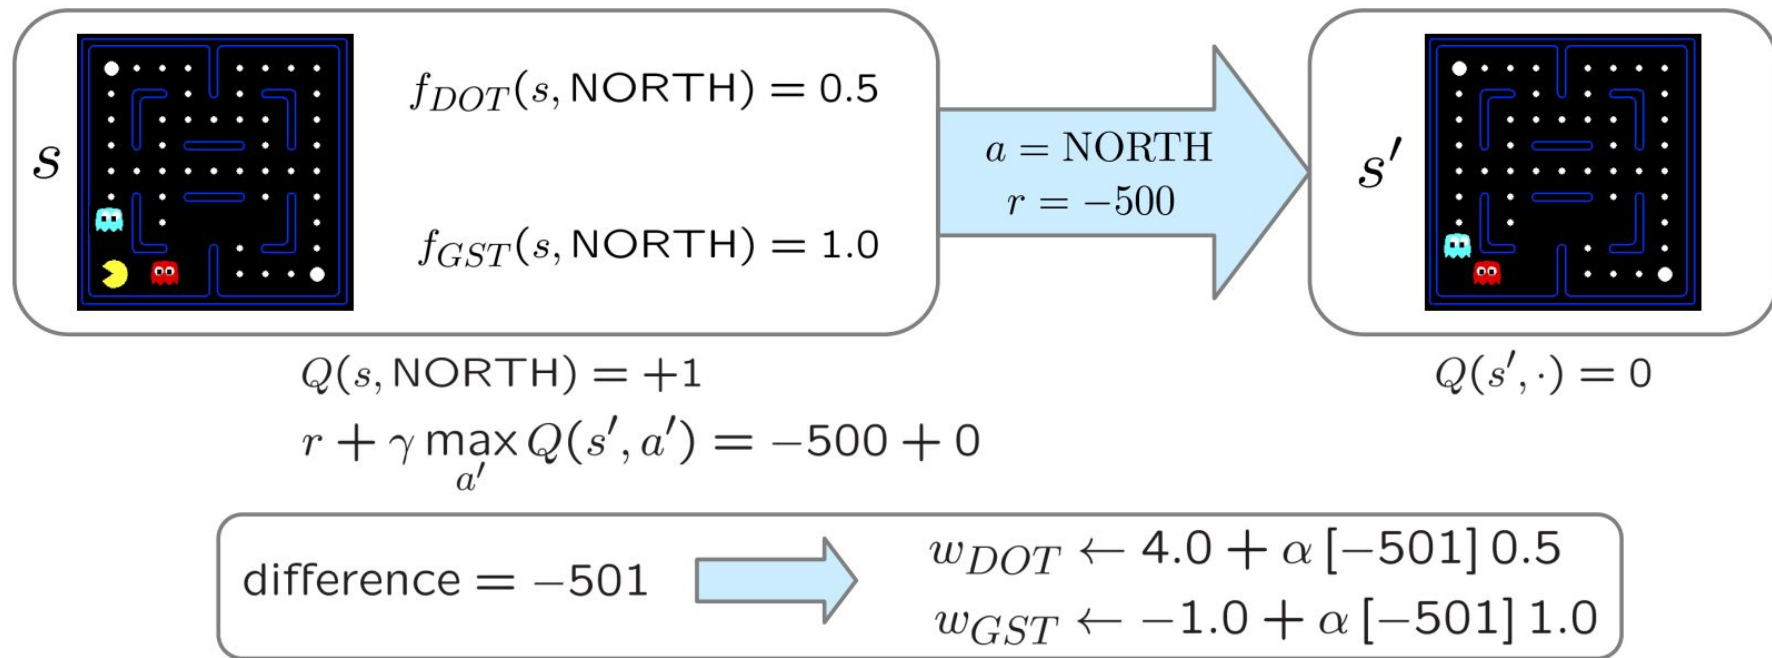
$$Q(s, a) \leftarrow Q(s, a) + \alpha [\text{difference}] \quad \text{Exact Q's}$$

$$w_i \leftarrow w_i + \alpha [\text{difference}] f_i(s, a) \quad \text{Approximate Q's}$$

- If something unexpectedly bad happens, blame the features that were on: disprefer all states with that state's features

Approximate Q-Learning


$$Q(s, a) = 4.0 f_{DOT}(s, a) - 1.0 f_{GST}(s, a)$$



$$Q(s, a) = 3.0 f_{DOT}(s, a) - 3.0 f_{GST}(s, a)$$


Least Squares Approximation

Linear Regression



Prediction:

$$\hat{y} = w_0 + w_1 f_1(x)$$




Prediction:

$$\hat{y}_i = w_0 + w_1 f_1(x) + w_2 f_2(x)$$

Linear Regression

$$\text{total error} = \sum_i (y_i - \hat{y}_i)^2 = \sum_i \left(y_i - \sum_k w_k f_k(x_i) \right)^2$$




Linear Regression

$$\text{error}(w) = \frac{1}{2} \left(y - \sum_k w_k f_k(x) \right)^2$$

$$\frac{\partial \text{error}(w)}{\partial w_m} = - \left(y - \sum_k w_k f_k(x) \right) f_m(x)$$

$$w_m \leftarrow w_m + \alpha \left(y - \sum_k w_k f_k(x) \right) f_m(x)$$



$$w_m \leftarrow w_m + \alpha \left[\underset{\text{“target”}}{r + \gamma \max_a Q(s', a')} - \underset{\text{“prediction”}}{Q(s, a)} \right] f_m(s, a)$$

Thank you